

Null-Model Choice for Rule-Defined Pattern Counts in Clustered Longitudinal Sequences:

A Case Study of Jewish Halachic Menstrual Anticipation Rules

Dvir Ross

ORCID: [0000-0001-8217-3359](https://orcid.org/0000-0001-8217-3359)

*Department of Software Engineering, Shenkar College of Engineering, Design and Art, Ramat
Gan, Israel*

Department of Computer Science, SCE — Shamoon College of Engineering, Beer Sheva, Israel

Author Note

Dvir Ross, Department of Software Engineering, Shenkar College of Engineering, Design and Art, Ramat Gan, Israel; Department of Computer Science, SCE — Shamoon College of Engineering, Beer Sheva, Israel.

The dataset analysed in this study was collected by [Fehring et al. \(2013a\)](#) and is publicly archived at Marquette University ([Fehring et al., 2013b](#)). The author declares no conflicts of interest. Analysis code, the filtered dataset and cached Monte Carlo replicates are available at <https://github.com/dvirross/VessetVectorTest>.

Article type: Case Study. Correspondence: dvirross@shenkar.ac.il or

RossDv@sce.ac.il

Abstract

Jewish religious law (*Halacha*) prescribes that a woman anticipate menstruation from recurring arithmetic patterns in her cycle history (the *vesset* rules). Each rule defines an exact event in a sequence of inter-menstrual intervals — three equal consecutive intervals, a constant first or second difference, a repeated weekly anchor — and the natural question is whether such events occur in real data more often than chance. The statistical difficulty is that the answer depends on what “chance” means for clustered longitudinal sequences whose rule-defined counts are dependent by construction. We analyse 1,554 cycles from 118 women under three exchangeability nulls for the same five-count vector: a global permutation null, a multinomial i.i.d. null, and a within-woman permutation null that conditions on each woman’s multiset of cycle lengths and randomises only their order. All are evaluated with two-sided Monte Carlo marginal tests and a joint quadratic-form (Mahalanobis) permutation statistic whose covariance is estimated from the null replicates. Under the global nulls the count vector is displaced far into the tail ($D_M = 5.08$, $p < .001$; Haflaga $z = 4.50$). Under the within-woman null it is not ($D_M = 2.93$, $p = .13$; smallest Holm-adjusted marginal $p = .19$), and all five observed counts lie at or below their conditional null means. A simulation study confirms nominal size for all tests, quantifies the power of the within-woman tests against autoregressive and exact-repetition alternatives, and shows the conclusions are insensitive to unequal follow-up. The contrast is consistent with between-woman heterogeneity in regularity generating the global excess, and provides no evidence that within-woman ordering produces additional rule-defined events; it does not establish the absence of serial dependence.

Keywords: exchangeability, Halacha, menstrual cycle, permutation test, within-subject null

1. Introduction

Menstrual cycle length varies both within and between women (Münster et al., 1992; Creinin et al., 2004; Harlow et al., 2000; Bull et al., 2019). The population distribution is bell-shaped with a pronounced right tail; within-woman consistency is moderate, with intraclass correlations typically between 0.3 and 0.5 (Dasharathy et al., 2012), and the menstrual cycle is increasingly treated as a within-person process requiring within-person modelling (Schmalenberger et al., 2021). Hierarchical state-space models exploit this structure for cycle-length prediction (Bortot et al., 2010), and recent multi-cohort work documents low-order autocorrelation in cycle lengths (Ecochard et al., 2024).

In Jewish religious law (*Halacha*) this regularity is also the object of a detailed legal system. A woman is required to anticipate the onset of her next menstruation from patterns in her prior cycle history; the anticipatory rules — the *halachic vessel rules* (*vessel*: menstruation) — govern abstinence obligations within *taharat hamishpacha*, the laws of family purity. Each rule defines an exact arithmetic event in the sequence of inter-menstrual intervals. Of the 17 rule types in the halachic literature, 10 require calendar or time-of-day data that standard cycle records do not contain; 7 are computable from interval lengths alone. Despite centuries of practical application and the availability of longitudinal cycle datasets, we are not aware of any formal randomisation-based test of whether these events occur more often than chance in empirical data (no systematic literature search was performed).

From a statistical point of view the question is a case of a general problem: testing whether counts of deterministically defined patterns in clustered longitudinal sequences exceed a null expectation, when (i) subjects contribute sequences of very unequal length (here 5–45 cycles), (ii) the counts are integer-valued and small, (iii) the counts are dependent by definition (one rule’s event set is a subset of another’s), and (iv) the same excess can be generated either by heterogeneity between subjects or by ordering within subjects. The choice of null is therefore not a technicality. Permutation inference for longitudinal and clustered data has been studied extensively (Friedrich et al., 2017; Lee and Braun, 2012), exchangeability-block permutation is standard in neuroimaging (Nichols and Holmes, 2002; Winkler et al., 2014), and conditional permutation tests that hold part of the data structure fixed have a modern theory (Berrett et al.,

2020; Hemerik and Goeman, 2018). This paper applies these established tools to the vessel counts and uses the application to make one point concrete: for the same data, a population-level exchangeability null and a within-subject conditional null return diametrically different answers, and only the latter addresses the question the rules implicitly pose.

We compare three nulls — a global permutation null H_G , a multinomial i.i.d. null H_{iid} and a within-woman permutation null H_W — using the same two-sided Monte Carlo marginal tests and the same joint quadratic-form statistic, a Mahalanobis distance from the simulated null mean with covariance estimated from the null replicates. The joint statistic is an adaptation of well-known quadratic-form permutation tests (Anderson, 2001; Pesarin and Salmaso, 2010; Székely and Rizzo, 2013); we make no claim of methodological novelty for it and instead verify its calibration by simulation. The study extends a line of quantitative halachic scholarship (Ross, 2022a, 2019, 2021, 2022b) and is presented as a Case Study.

2. Background

2.1 The Polynomial Hierarchy of Vessel Rules

Let $L_1, \dots, L_n$ denote one woman's cycle lengths (days) and define the *haflaga* $H_k = L_k + 1$, the interval from the first day of menstruation k to the first day of menstruation $k + 1$, following halachic convention (Ross, 2021). The five rule types analysed here are:

Haflaga. $H_k = H_{k-1} = H_{k-2}$: three or more consecutive equal haflaga values. Degree-0 (constant) sequence; minimum run 3 cycles.

Dilug. $H_k - H_{k-1} = H_{k-1} - H_{k-2} = d \neq 0$: constant non-zero first difference. Degree-1 (linear); minimum run 3.

Dilug-in-Dilug. Constant non-zero second difference, $(H_k - H_{k-1}) - (H_{k-1} - H_{k-2}) = (H_{k-1} - H_{k-2}) - (H_{k-2} - H_{k-3}) = d \neq 0$. Degree-2 (quadratic); minimum run 4.

Week. $H_k = H_{k-1} \in \mathcal{W}$, where $\mathcal{W} = \{H_{\min} + 3, H_{\min} + 10, \dots\}$ are values spaced 7 apart (same weekday at each onset). Degree-0 anchored to weekly values; minimum run 2.

Week-Dilug. $H_k = H_{k-1} = 30$. A haflaga of 30 (29 days, four weeks plus one day) is the unique interval for which each successive onset falls one weekday later, producing an

arithmetic progression of weekdays — hence the name (Hebrew *shavu'a be-dilug*; Ross, 2022b). Degree-0 at the canonical anchor; minimum run 2.

Haflaga, Dilug and Dilug-in-Dilug follow increasing polynomial order in cycle index; Week and Week-Dilug are specialisations of the degree-0 case anchored to calendar values. Two further computable rules (*Dilug Chozar Chalila* and *Haflaga Chozar Chalila*, repeating multi-value cycles of intervals) require windows of at least six to nine cycles; they are discussed in Section 3.1.

2.2 Dependence Structure Implied by the Definitions

Two relationships follow from the definitions and are relevant to how the counts should be tested jointly.

Haflaga-30 $\subseteq$ Week-Dilug (deterministic containment). — Week-Dilug requires two consecutive haflaga values equal to 30; a Haflaga sequence with value 30 requires three. Every Haflaga-30 sequence therefore generates a Week-Dilug event in its first two cycles. This is set-theoretic inclusion, not a stochastic association: in the data, the 3 Haflaga sequences with $H = 30$ contribute directly to the Week-Dilug count of 18.

Haflaga and Dilug-in-Dilug (algebraic proximity). — When the second difference $|d|$ of a Dilug-in-Dilug sequence is 1, its first differences change by one day per cycle and the haflaga sequence is nearly constant. In the data, 16 of 27 Dilug-in-Dilug sequences have $|d| = 1$. This makes positive co-occurrence with Haflaga plausible but does not entail it.

A third relation, mutual exclusion of Dilug ($\Delta H \neq 0$) and Week-Dilug ($\Delta H = 0$) within the same window, implies a weak negative tendency for that pair. The five counts are thus non-independent by construction, which motivates estimating their joint null covariance rather than treating them as separate tests.

2.3 Prior Work and Dataset Overlap

Statistical characterisation of cycle-length variability is extensive (Münster et al., 1992; Creinin et al., 2004; Harlow et al., 2000; Fehring et al., 2006; Dasharathy et al., 2012; Symul et al., 2019; Bull et al., 2019), and hierarchical models for sequential prediction exist (Bortot et al.,

2010). Quantitative halachic analyses have modelled cycle transitions with Markov chains (Ross, 2022a), examined the reliability of self-reported cycle data in a halachic setting (Ross, 2019), and analysed abstinence rules statistically; none tested pattern frequencies against a randomisation null.

Ecochard et al. (2024) report evidence of an internal circamonthly oscillator using multi-cohort data that include a North American cohort collected by Fehring and Schneider, two of the investigators of the dataset analysed here (Fehring et al., 2013a). Their Marquette cohort (1,649 cycles, 159 women) and the present sample (1,554 cycles, 118 women after exclusions) are probably drawn from the same data-collection infrastructure and should not be regarded as independent samples. They document second- to fourth-order autocorrelation in cycle lengths. We note the implication for H_W precisely: within-woman permutation preserves each woman’s multiset of cycle lengths and her sequence length, and *destroys* the observed ordering and therefore any serial dependence. Autocorrelation is not preserved under H_W ; rather, H_W is the reference distribution in which it is absent, so a rejection of H_W would indicate ordering structure that produces rule-defined events, and a non-rejection indicates only that these specific count statistics are not unusual relative to random reorderings of each woman’s own values.

3. Methods

3.1 Dataset

We analysed the publicly archived longitudinal dataset of Fehring et al. (2013a), collected in a randomised trial of two internet-supported natural family planning (NFP) methods (Fehring et al., 2013b). The archived filtered file retains cycles of 18–54 days with complete cycle-length records (Creinin et al., 2004); the analysed sample comprises **1,554 cycles from 118 women** ($M = 13.2$ cycles per woman, $SD = 6.1$, range 5–45). Cycle lengths ranged from 18 to 54 days ($M = 29.3$, $SD = 3.9$, $Mdn = 28.5$, skewness = 1.29); 9.6% of cycles were 35 days or longer. Women were aged 21–43 years ($M = 30.6$, $SD = 5.7$); parity was not recorded. The data come from two charting cohorts ($n = 69$ and 49 women in the filtered sample); cohort membership was not used in the analysis. All lengths were incremented by 1 to the haflaga convention, giving $M = 30.3$, $SD = 3.9$.

Because the pattern detectors operate on *consecutive* cycles, it matters whether the cycle exclusions could have created artificial adjacency across a removed cycle. In the filtered file every woman’s recorded `CycleNumber` runs $1, 2, \dots, n_i$ without gaps (verified programmatically and asserted in the analysis code), so no pattern window spans an excluded cycle; the exclusions removed whole women or terminal cycles only. The minimum of 5 cycles per woman is the observed minimum after filtering; no additional minimum was imposed. The script that produced the filtered file from the raw archive was not preserved and is not part of the code release; the filtered file itself is archived with the code.

Of the seven rules computable from cycle lengths, the two repeating-cycle rules (*Dilug Chozer Chalila*, *Haflaga Chozer Chalila*) had zero occurrences, and the analysed family of five was fixed after this was observed. We state this plainly: the family is data-informed, not pre-specified. Its consequences are limited — a zero-count component cannot be included in a Mahalanobis statistic (its null covariance is degenerate at the observed sequence lengths), and adding two marginal tests with observed count 0 would not alter any conclusion below — but the family should be read as primary rather than confirmatory.

As a descriptive characterisation, the 20 women with at least one Haflaga event had a mean within-woman SD of cycle length of 1.54 days (between-woman SD 0.59), compared with 3.17 days (1.75) for the 35 women with no event of any type (Mann–Whitney $U = 133$, $p < .001$; Welch $t = -5.05$, $df = 45.5$; 95% CI for the difference $[-2.29, -0.98]$ days; Cohen’s $d = -1.13$). Because Haflaga is *defined* by repeated equal intervals, a lower SD among these women is expected by construction; this comparison describes the women the rule selects and is not independent validation of the rule. Within-woman SD estimates also depend on the number of cycles observed (5–45). Figure 1 shows the distribution and per-woman variability; Figure 2 summarises the data flow.

3.2 Pattern Detection

Counting is performed within each woman’s sequence; sequences are never concatenated across women, and the woman boundaries (s_i, f_i) are enforced in every counting loop. Each function counts *distinct maximal contiguous runs* satisfying its rule: a run longer than the minimum counts once (four equal haflaga values are one Haflaga event), and a qualifying run that resumes




Figure 1: Cycle length distribution (filtered dataset). *Left:* histogram and kernel density estimate of L_k with mean (dashed) and median (dotted); skewness = 1.29. *Right:* per-woman mean vs. within-woman SD, shaded by cycle count.

after a discordant cycle counts again. To fix the convention: $(29, 29)$ qualifies as Week (if $29 \in \mathcal{W}$) but not as Haflaga; $(28, 28, 28)$ is a Haflaga; $(27, 29, 31, 33)$ is a Dilug ($d = 2$) with second differences $(0, 0)$ and hence not a Dilug-in-Dilug; $(27, 28, 30, 33)$ has first differences $(1, 2, 3)$ and second differences $(1, 1)$ and is a Dilug-in-Dilug.

Two implementations were written — a reference loop version and a vectorised version used for simulation — and checked to agree on the observed data and on 9,000 random null sequences. The observed count vector is $\mathbf{O} = (26, 61, 32, 18, 27)$ for (Haflaga, Dilug, Week, Week-Dilug, Dilug-in-Dilug). The code, the filtered data, and all Monte Carlo replicates are available at <https://github.com/dvirross/VesetVectorTest>.

3.3 Null Models

Three nulls encode three scientific questions (Table 1).

H_G (global). All 1,554 haflaga values are exchangeable across women: the pooled values are permuted without replacement and re-assigned to women preserving each n_i . Tests whether the counts exceed what a homogeneous population would produce.

H_{iid} (multinomial). Each woman's n_i values are drawn independently from the pooled empirical distribution $\hat{F}$ (Ross, 2022b). Stronger than H_G in that observed values, not only their ordering, are replaced.




Figure 2: Participant and data flow. Source dataset (Fehring et al., 2013a); inclusion criteria (cycle length 18–54 days; complete records); analysed sample of 118 women and 1,554 cycles; five pattern types; three null models, each with $B = 50,000$ replicates.

H_W (within-woman). Conditional on each woman’s observed multiset $\{H_{i,1}, \dots, H_{i,n_i}\}$ and on n_i , all $n_i!$ orderings are equally likely. Each woman’s sequence is permuted uniformly at random without replacement (indices are permuted, so ties are handled exactly). This is a conditional Monte Carlo test (Hemerik and Goeman, 2018): it preserves each woman’s level of regularity and cycle count, destroys within-woman ordering, and tests whether the ordering generates rule-defined events beyond what her own values in random order would. It is not a general test of serial independence: an autocorrelated or trending process that does not produce additional exact events will not reject it.

Table 1: The three null models.

	H_G (global permutation)	H_{iid} (multinomial)	H_W (within-woman permutation)
Randomised	assignment of pooled cycle values to positions	each value, drawn i.i.d. from $\hat{F}$	order of each woman’s own values
Preserved	pooled multiset; each n_i	each n_i ; pooled proportions	each woman’s multiset and n_i
Destroyed	woman identity; ordering	all dependence; heterogeneity	within-woman ordering only
Scientific null	counts arise from a homogeneous population	counts arise from independent draws	counts arise from random ordering of each woman’s values

A within-woman *bootstrap* (sampling each woman’s values with replacement) was considered and rejected as the conditional null: with-replacement draws are independent, whereas permutation draws are negatively dependent (a value used at one position is less available at the next), so the bootstrap inflates the expected count of repetition-based patterns (Haflaga null mean 54.5 vs. 31.8; verified with $B = 50,000$). The two procedures encode different nulls — i.i.d. resampling versus exchangeability of the observed values — and are not more or less conservative versions of one another; permutation is the exact conditional test for the stated exchangeability null.

3.4 Statistical Inference

For pattern j and null m , with $B = 50,000$ replicate count vectors $\mathbf{S}^{(m,r)}$, we report the null mean $\bar{S}_j^{(m)}$ and SD $\sigma_j^{(m)}$, the standardised departure $z_j^{(m)} = (O_j - \bar{S}_j^{(m)})/\sigma_j^{(m)}$, and the two-sided

Monte Carlo p-value

$$p_j^{(m)} = \min \left\{ 1, 2 \min \left(\frac{b_j^+ + 1}{B+1}, \frac{b_j^- + 1}{B+1} \right) \right\}, \quad (1)$$

where $b_j^+ = \#\{r : S_j^{(m,r)} \geq O_j\}$ and $b_j^- = \#\{r : S_j^{(m,r)} \leq O_j\}$, following the $(b+1)/(B+1)$ convention (Davison and Hinkley, 1997; Phipson and Smyth, 2010). The same convention is used everywhere (tables, figures, algorithm). Two-sided tests are used because departures in either direction are informative for these rules. Across the five marginal tests the family-wise error rate is controlled by Holm’s step-down procedure, which is valid under arbitrary dependence.

The **joint test** uses the Mahalanobis distance of $\mathbf{O}$ from the simulated null mean,

$$D_M^{(m)} = \sqrt{(\mathbf{O} - \hat{\boldsymbol{\mu}}^{(m)})^\top (\hat{\boldsymbol{\Sigma}}^{(m)})^{-1} (\mathbf{O} - \hat{\boldsymbol{\mu}}^{(m)})}, \quad (2)$$

where $\hat{\boldsymbol{\mu}}^{(m)}$ and $\hat{\boldsymbol{\Sigma}}^{(m)} = (B-1)^{-1} \sum_r (\mathbf{S}^{(m,r)} - \bar{\mathbf{S}}^{(m)})(\mathbf{S}^{(m,r)} - \bar{\mathbf{S}}^{(m)})^\top$ are the mean and sample covariance of the null replicates (not of the observations). The p-value is the fraction of replicates whose own D_M , computed with the same $\hat{\boldsymbol{\mu}}$ and $\hat{\boldsymbol{\Sigma}}$, is at least the observed value, again with the $(b+1)/(B+1)$ convention. This is a single global test of the multivariate null and is inherently two-sided; it is a standard permutation quadratic form (Anderson, 2001; Pesarin and Salmaso, 2010) closely related to Hotelling’s T^2 with the parametric reference distribution replaced by a Monte Carlo one, and to energy-distance tests (Székely and Rizzo, 2013). Because the same replicates estimate $\hat{\boldsymbol{\Sigma}}$ and form the reference distribution, we also report a **split-batch** version ($\hat{\boldsymbol{\mu}}, \hat{\boldsymbol{\Sigma}}$ from the first 25,000 replicates; reference distribution from the last 25,000) and verify calibration by simulation (Section 3.6). Condition numbers of $\hat{\boldsymbol{\Sigma}}$ (4.4–5.4) indicate a numerically stable inversion; they say nothing about power. The joint test and the Holm family answer different questions (one global null versus family-wise error control over five component nulls) and are reported as complements.

The joint test on all five patterns is the primary multivariate analysis. An **exploratory** joint test on the three structurally linked patterns (Haflaga, Week-Dilug, Dilug-in-Dilug; Section 2.2) is reported to visualise the displacement geometry in three dimensions (Figure 8); the restriction was not pre-registered and the three-pattern result is not used for inference.

3.5 Pairwise Dependence Analysis (Exploratory)

The joint test treats the five counts as one vector; the pairwise analysis asks a different and secondary question — whether women who produce one pattern type tend to produce another — as a consistency check against the relationships derived in Section 2.2. Each design choice below follows from a specific feature of the data.

Unit of analysis: the woman. — Pattern events are defined within a woman’s sequence and cycles within a woman are not independent, so the woman is the only exchangeable unit for an association test. The analysis therefore works with 118 per-woman summaries, not with 1,554 cycles.

Rates rather than counts. — A woman with 45 cycles has roughly nine times as many opportunities to produce *any* pattern as a woman with 5 cycles, so raw per-woman counts of two patterns are positively correlated through follow-up length alone. Dividing each count by the number of cycles removes the first-order part of this exposure effect. It does not remove it entirely — a rate from 5 cycles is far noisier than one from 45, and the minimum run lengths (2, 3 or 4 cycles) mean short sequences have proportionally fewer eligible windows — which is one reason the analysis is labelled exploratory.

Spearman rather than Pearson correlation. — Per-woman rates are bounded at zero, heavily zero-inflated (most women have no event of a given type) and right-skewed. Pearson’s coefficient would be dominated by the few women with several events; the rank correlation measures monotone association and is insensitive to the magnitude of those few values.

Permutation rather than asymptotic p-values. — With 118 women, many tied zeros and skewed rates, the usual t -approximation for the null distribution of r_s is unreliable. Instead one rate vector is permuted across women 50,000 times; this breaks any association between the two patterns while preserving both marginal distributions exactly, giving a two-sided Monte Carlo p-value (proportion of permutations with $|r_s^*| \geq |r_s|$, with the $(b + 1)/(B + 1)$ convention).

Holm correction over ten pairs. — Five patterns yield ten pairs, of which only two carry a structural expectation. Holm’s step-down procedure controls the family-wise error rate at $\alpha = .05$ over all ten and is valid under any dependence among the tests, which matters here

because the pairs share patterns.

Binary robustness check. — Because the rate metric depends on the exposure normalisation and on the few women with several events, the analysis is repeated with each woman coded 1/0 for ever having established each pattern. This discards magnitude entirely. Each pair is then a 2×2 table tested with Fisher’s exact test (conditional on the margins) and with a permutation χ^2 statistic (permuting one indicator across women), again Holm-corrected. Agreement between the rate-based and binary results indicates that the conclusions do not depend on the choice of metric.

What a positive association can and cannot mean. — For Haflaga $\times$ Week-Dilug a positive association is partly guaranteed: every Haflaga sequence at $H = 30$ contains a Week-Dilug event, so the test partly re-discovers the definitions. For Haflaga $\times$ Dilug-in-Dilug the definitions do not entail co-occurrence, so an association is an empirical observation about the women in this sample; it is compatible with the proximity argument but would require a generative model to attribute to it. In neither case does the pairwise result bear on the primary question of whether ordering produces excess events.

3.6 Simulation Validation and Sensitivity Analyses

Three questions were addressed by simulation, all using the counting code and test procedures of the main analysis.

Size. — For each null, fresh pseudo-observed datasets (10,000) were drawn from the null and analysed against an independent reference batch of 20,000 replicates. We report the rejection rate at $\alpha = .05$ of each two-sided marginal test, of the Holm family (any rejection), and of the joint test in its plug-in form (reference D_M distribution and $\hat{\Sigma}$ from the same batch, as in the main analysis) and split-batch form.

Power of the within-woman tests. — Datasets with genuine within-woman temporal structure were generated in two families and each analysed with its own 2,000 within-woman permutations (400 datasets per scenario). *AR(1)*: each woman’s sequence is a Gaussian first-order autoregressive process with her observed mean and SD and coefficient $\rho \in \{0, .25, .5, .75\}$, rounded to whole days and clipped to the observed range. *Persistence*: each cycle exactly

repeats the previous haflaga with probability $q \in \{0, .1, .2, .3\}$ and is otherwise drawn from the woman's own observed values; this is the mechanism the Haflaga rule presupposes. In both families the zero level is exchangeable within woman and serves as a further size check under heterogeneous, realistic marginals.

Sensitivity to unequal follow-up. — The H_W analysis was repeated with every woman truncated to her first 12 recorded cycles (20,000 replicates), and a leave-one-woman-out analysis recomputed the H_W joint and marginal p-values 118 times using per-woman replicate counts.

Algorithm 1 Simulation pipeline

Require: Sequences $\{(H_{i,1}, \dots, H_{i,n_i})\}_{i=1}^N$, patterns $j = 1, \dots, 5$, replicates B , seed

- 1: $\mathbf{O} \leftarrow \text{Count}(\mathbf{H})$
- 2: **for** $r = 1$ to B **do**
- 3: $\mathbf{S}^{(G,r)} \leftarrow \text{Count}(\text{Permute}(\bigcup_i \{H_{i,k}\}))$ ▷ pooled values re-assigned preserving n_i
- 4: $\mathbf{S}^{(M,r)} \leftarrow \text{Count}(\text{i.i.d. draws from } \hat{F}, n_i \text{ per woman})$
- 5: **for** $i = 1$ to N **do** $\mathbf{H}_i^{(W)} \leftarrow \text{Permute}(H_{i,1}, \dots, H_{i,n_i})$
- 6: **end for**
- 7: $\mathbf{S}^{(W,r)} \leftarrow \text{Count}(\mathbf{H}^{(W)})$
- 8: **end for**
- 9: **for** $m \in \{G, M, W\}$ **do**
- 10: $\hat{\boldsymbol{\mu}}^{(m)} \leftarrow B^{-1} \sum_r \mathbf{S}^{(m,r)}$; $\hat{\Sigma}^{(m)} \leftarrow \text{SampleCov}(\{\mathbf{S}^{(m,r)}\})$
- 11: **for** $j = 1$ to 5 **do**
- 12: $b_j^+ \leftarrow |\{r : S_j^{(m,r)} \geq O_j\}|$; $b_j^- \leftarrow |\{r : S_j^{(m,r)} \leq O_j\}|$
- 13: $p_j^{(m)} \leftarrow \min\{1, 2 \min((b_j^+ + 1)/(B + 1), (b_j^- + 1)/(B + 1))\}$ ▷ two-sided
- 14: **end for**
- 15: $D_M^{(m)} \leftarrow \sqrt{(\mathbf{O} - \hat{\boldsymbol{\mu}}^{(m)})^\top (\hat{\Sigma}^{(m)})^{-1} (\mathbf{O} - \hat{\boldsymbol{\mu}}^{(m)})}$
- 16: $D_M^{(m,r)} \leftarrow \text{same with } \mathbf{S}^{(m,r)} \text{ in place of } \mathbf{O}, \text{ for all } r$
- 17: $p_{D_M}^{(m)} \leftarrow (|\{r : D_M^{(m,r)} \geq D_M^{(m)}\}| + 1)/(B + 1)$
- 18: **end for**

4. Results

4.1 Global Nulls: Marginal Tests

Table 2 gives the marginal results under H_G and H_{iid} ; Figures 3 and 4 show the null distributions and standardised departures. Haflaga is far in the upper tail: 26 observed against a null mean of 11.31 ($z = 4.50$); only 4 of 50,000 H_G replicates reached 26 (two-sided $p = .0002$; Holm-adjusted $p = .001$). No other pattern is significant. Dilug ($z = 1.84$, $p = .085$, Holm $p = .34$) is at most hypothesis-generating; Week, Week-Dilug and Dilug-in-Dilug lie within their null

ranges. The two global nulls agree closely (Figure 5; largest difference in two-sided p-values 0.11, for Week, whose null SD is larger under H_{iid}).

Table 2: Marginal results under the global nulls ($B = 50,000$). Null mean and SD are from H_G ; z is the standardised departure under H_G ; p-values are two-sided (Eq. 1); Holm adjustment is across the five patterns within each null. Bold: Holm-adjusted $p < .05$.

Pattern	Obs.	Null mean	Null SD	z	p (H_G)	Holm (H_G)	p (H_{iid})
Haflaga	26	11.31	3.27	+4.50	< .001	.001	< .001
Dilug	61	48.92	6.57	+1.84	.085	.340	.082
Week	32	27.02	4.01	+1.24	.265	.794	.373
Week-Dilug	18	16.25	3.25	+0.54	.688	1.00	.743
Dilug-in-Dilug	27	27.17	5.00	−0.03	1.00	1.00	1.00




Figure 3: Empirical null probability mass functions under H_G (left bar of each pair) and H_{iid} (right bar) with the observed count (dashed line). Saturated bars mark the tail at or beyond the observed count on the observed side; the two-sided p-value is twice the smaller tail probability (Eq. 1) and is printed in each panel for both nulls.

4.2 Within-Woman Null

The global nulls remove between-woman heterogeneity in regularity: they ask whether the counts exceed what a homogeneous population would produce, not whether each woman’s *ordering* produces events beyond what her own values predict. H_W asks the latter. Table 3 and




Figure 4: Standardised departures z_j under H_G (solid) and H_{iiid} (hatched). Dashed and dotted lines mark $|z| = 1.96$ and 2.576 . The two nulls give different z for the same count because each has its own mean and SD (e.g. Week: 1.24 vs. 0.98).




Figure 5: Empirical CDFs of the H_G (solid) and H_{iiid} (dashed) null distributions; dotted vertical line at the observed count.

Figure 6 give the results.

Table 3: Within-woman permutation null H_W ($B = 50,000$; all 118 women; each woman’s sequence permuted without replacement). p-values are two-sided (Eq. 1); Holm adjustment across the five patterns. LOO range: range of the two-sided p-value over the 118 leave-one-woman-out analyses.

Pattern	Obs.	Null mean	Null SD	z	p	Holm	LOO range
Haflaga	26	31.84	4.68	−1.25	.254	1.00	.12–.39
Dilug	61	67.04	7.38	−0.82	.460	1.00	.34–.60
Week	32	33.79	3.88	−0.46	.749	1.00	.53–.93
Week-Dilug	18	20.91	3.22	−0.90	.460	1.00	.31–.69
Dilug-in-Dilug	27	39.12	5.84	−2.07	.038	.192	.014–.073

Joint test (five patterns): $D_M = 2.93$, $p = .126$ (6,285 of 50,000 replicates at or beyond the observed value); split-batch $D_M = 2.95$, $p = .119$; leave-one-woman-out range of p : .055–.224.

Under H_W no pattern is significant after Holm correction, and all five observed counts lie *below* their conditional null means. The most extreme component is Dilug-in-Dilug ($z = -2.07$; nominal two-sided $p = .038$, Holm $p = .19$): fewer second-order progressions were observed than random reorderings of the same values produce. The direction is noted; the mechanism is not identified by these data, and the result is exploratory. The joint test is non-significant ($D_M = 2.93$, $p = .13$) and essentially unchanged in the split-batch form ($D_M = 2.95$, $p = .12$).

The shift between nulls is large: for Haflaga the H_G null is centred at 11.3 and the H_W null at 31.8, above the observed 26. The reason is that women with narrow cycle-length multisets generate many exact repetitions under *any* ordering of their values; conditioning on the multiset absorbs this.

4.3 Joint Tests

Under H_G the five-pattern Mahalanobis distance is $D_M = 5.08$ ($p < .001$; 6 of 50,000 replicates at or beyond) and under H_{iid} it is $D_M = 4.87$ ($p < .001$; 17 of 50,000); the split-batch values are 5.07 and 4.87 with the same p-values. Under H_W , $D_M = 2.93$, $p = .126$ (Table 3; Figure 7). The displacement under the global nulls is driven by Haflaga, which has the smallest null SD (3.27) and the largest standardised departure and therefore receives the largest weight in the quadratic form; the null covariance matrices are well conditioned (condition numbers 4.6, 4.4




Figure 6: Smoothed null densities (for visual comparison only; inference uses the empirical PMFs) under H_G (solid), H_{iid} (dashed) and H_W (dash-dotted) for each pattern, with the observed count (dotted vertical line) and the two-sided H_W p -value in each panel.

and 5.4 for H_G , H_{iid} , H_W) with all pairwise null correlations below 0.17 in absolute value.

The exploratory three-pattern test (Haflaga, Week-Dilug, Dilug-in-Dilug) gives $D_M = 4.50$ ($p < .001$), 4.39 ($p < .001$) and 2.68 ($p = .067$) under H_G , H_{iid} and H_W respectively, consistent with the five-pattern analysis (Figure 8).




Figure 7: Null distributions of the five-pattern Mahalanobis distance D_M (Eq. 2) under H_G (left), H_{iid} (centre) and H_W (right); dashed line at the observed value.




Figure 8: Exploratory three-pattern subspace: 3,000 of the 50,000 null replicates (points) and the observed vector (26, 18, 27) in the (Haflaga, Week-Dilug, Dilug-in-Dilug) space under H_G (top left), H_{iid} (top right) and H_W (bottom), so that all five counts are visible at once: marker size encodes the Dilug count (size legend in each panel) and marker shade the Week count (colour bar). The filled circle is the null mean; the dashed segment is the displacement to the observed vector, with its Euclidean length D_E and Mahalanobis length D_M in the legend. Primary inference uses all five patterns (Table 3).

4.4 Pairwise Dependence (Exploratory)

The methods and their rationale are given in Section 3.5. Section 2.2 leads to two expectations: a positive Haflaga $\times$ Week-Dilug association, which is partly an arithmetic consequence of containment rather than an empirical hypothesis, and a plausible but not entailed positive Haflaga $\times$ Dilug-in-Dilug association from polynomial proximity. No expectation is held for the other eight pairs, so the analysis is a consistency check, not a search.

Table 4 and Figure 9 give the rate-based Spearman correlations. Haflaga $\times$ Dilug-in-Dilug is positive and survives Holm correction ($r_s = 0.27$, $p_{\text{perm}} = .004$, Holm $p = .038$): women with a higher rate of exact repetitions also tend to have a higher rate of second-order progressions, most of which ($|d| = 1$ in 16 of 27) are near-constant sequences. Haflaga $\times$ Week-Dilug is positive, as containment requires, but does not survive correction ($r_s = 0.23$, $p_{\text{perm}} = .011$, Holm $p = .10$); only 3 of the 26 Haflaga sequences are at $H = 30$, so the guaranteed part of the association is small. The remaining eight pairs are null (all Holm $p = 1$), including Dilug $\times$ Week-Dilug, whose weak negative sign ($r_s = -0.12$) is in the direction implied by their mutual exclusion but is not distinguishable from zero.

The binary (ever/never) version reproduces the ordering and the two smallest p-values (Haflaga $\times$ Dilug-in-Dilug: $\chi^2 = 9.99$, $p_{\text{perm}} = .004$, Holm $p = .042$, Fisher $p = .004$; Haflaga $\times$ Week-Dilug: $\chi^2 = 7.26$, $p_{\text{perm}} = .013$, Holm $p = .12$), so the findings do not depend on the rate normalisation or on the few women with multiple events. Three cautions apply. The associations are between women, not within sequences, and say nothing about temporal ordering. They remain exposed to the residual exposure difference between women with 5 and 45 cycles (Section 3.5). And the surviving association is compatible with the proximity mechanism but is not evidence for it: attributing co-occurrence to a mechanism would require an explicit generative model for the sequences, which this paper does not fit.

4.5 Simulation Validation and Sensitivity

Size. — Table 5 reports empirical rejection rates under each null for 10,000 fresh pseudo-observed datasets (Monte Carlo standard error about 0.002 at $\alpha = .05$). Every procedure holds its nominal level: the two-sided marginal tests reject between 2.9% and 4.8% of the time (below

Table 4: Pairwise Spearman correlations between per-woman pattern rates ($N = 118$), sorted by p_{perm} (two-sided, 50,000 permutations); p_{Holm} across the 10 pairs. §: pair with an analytically expected positive association (Section 2.2). Bold: $p_{\text{Holm}} < .05$.

Pair	r_s	p_{perm}	p_{Holm}
Haflaga $\times$ Dilug-in-Dilug§	+0.27	.004	.038
Haflaga $\times$ Week-Dilug§	+0.23	.011	.102
Dilug $\times$ Week-Dilug	−0.12	.206	1.00
Haflaga $\times$ Week	+0.10	.277	1.00
Week $\times$ Dilug-in-Dilug	+0.09	.326	1.00
Dilug $\times$ Week	+0.09	.337	1.00
Week $\times$ Week-Dilug	+0.06	.476	1.00
Week-Dilug $\times$ Dilug-in-Dilug	+0.07	.480	1.00
Haflaga $\times$ Dilug	+0.04	.668	1.00
Dilug $\times$ Dilug-in-Dilug	−0.02	.819	1.00




Figure 9: Spearman correlations between per-woman pattern rates. §: analytically expected positive pair; *: Holm-adjusted $p < .05$. All pairwise tests are exploratory.

nominal because the counts are discrete), the Holm family between 3.2% and 3.7%, and the joint test between 4.8% and 5.3% at $\alpha = .05$ and between 0.7% and 1.1% at $\alpha = .01$. The plug-in and split-batch versions of the joint test are indistinguishable, so re-using the replicates for $\hat{\Sigma}$ has no detectable effect at $B \geq 20,000$.

Table 5: Empirical size at nominal $\alpha = .05$ (10,000 pseudo-observed datasets per null, each tested against an independent batch of 20,000 replicates). Joint test shown in plug-in form ($\hat{\Sigma}$ and reference distribution from the same batch) and split-batch form; values at $\alpha = .01$ in parentheses.

Null	Marginal two-sided tests					Holm family	Joint D_M	
	Haf.	Dilug	Week	W-D	DiD		plug-in	split-batch
H_G	.029	.047	.034	.031	.046	.033	.048 (.011)	.052 (.010)
H_{iid}	.048	.038	.040	.038	.046	.032	.049 (.007)	.047 (.007)
H_W	.043	.041	.043	.042	.040	.037	.053 (.009)	.053 (.010)

Power of the within-woman tests. — Table 6 and Figure 10 give rejection rates of the H_W tests for datasets with genuine within-woman ordering (400 datasets per scenario; Monte Carlo $SE \leq 0.025$). At $\rho = 0$ and $q = 0$, which are exchangeable within woman, all tests reject at nominal rates under these heterogeneous marginals (joint 6.2%, SE 1.2%). Autoregressive dependence with $\rho = 0.25$ — which raises the expected Haflaga count by about five events over its conditional null mean — is detected by the joint test in 51% of datasets and by the Holm family in 26%; at $\rho = 0.5$ the joint test rejects in 99.5% and at $\rho = 0.75$ always. Exact-repetition persistence is detected far more readily: when one cycle in ten copies its predecessor ($q = 0.1$) the joint test rejects in 98% of datasets and the Haflaga marginal test in 96%; at $q \geq 0.2$ detection is certain. The Dilug-in-Dilug marginal has no power against persistence (repetition suppresses second-order progressions), and Week-Dilug has the least power against autoregression, as expected from their definitions.

These results calibrate the main finding. Had within-woman ordering produced an excess of roughly five Haflaga events in this sample, the joint test would have detected it about half the time; an excess of 14 or more would almost certainly have been detected. The observed Haflaga count is six *below* its conditional null mean. Serial dependence weaker than $\rho \approx 0.25$, or dependence that does not manifest as whole-day autoregression or exact repetition, may go undetected by these statistics.

Table 6: Power of the H_W tests at $\alpha = .05$ against within-woman temporal-dependence alternatives (400 datasets per scenario, each analysed with its own 2,000 within-woman permutations). $\Delta\text{Haf.}$: mean observed Haflaga count minus its mean conditional null expectation. The $\rho = 0$ and $q = 0$ rows are exchangeable within woman and serve as size checks.

Alternative	Level	$\Delta\text{Haf.}$	Haf.	Dilug	Week	W-D	DiD	Holm	Joint
AR(1) ρ	0	-0.1	.04	.05	.04	.04	.04	.03	.06
	0.25	+5.0	.19	.24	.10	.06	.18	.26	.51
	0.50	+13.8	.76	.74	.35	.31	.44	.89	.995
	0.75	+33.3	1.00	.97	.78	.67	.73	1.00	1.00
Persistence q	0	+0.2	.04	.04	.04	.04	.06	.05	.06
	0.1	+24.2	.96	.19	.61	.41	.04	.94	.98
	0.2	+52.7	1.00	.60	.96	.86	.04	1.00	1.00
	0.3	+80.4	1.00	.93	.99	.96	.17	1.00	1.00




Figure 10: Rejection rates of the H_W tests at $\alpha = .05$ against within-woman AR(1) alternatives (left) and exact-repetition alternatives (right); the leftmost point of each panel is a size check. Solid: joint five-pattern D_M ; dashed: any Holm-corrected marginal; dotted: selected marginal tests. Dash-dotted horizontal line at 0.05.

Sensitivity to unequal follow-up. — Eleven women with 20 or more cycles contribute 20% of cycles and 35% of Haflaga events (21%, 28%, 11% and 30% of Dilug, Week, Week-Dilug and Dilug-in-Dilug events); the Spearman correlation between a woman’s cycle count and her Haflaga events is 0.28. When every woman is truncated to her first 12 cycles (1,289 cycles), the observed counts are (19, 51, 21, 16, 21), all z remain negative (−1.18, −0.32, −0.89, −0.68, −1.78), no marginal test approaches significance (smallest two-sided $p = .080$, for Dilug-in-Dilug), and the joint test gives $D_M = 2.49$, $p = .29$. Leaving out any single woman moves the joint H_W p-value within .055–.224 and never below .05 (Table 3); the nominal Dilug-in-Dilug marginal stays below .05 in 114 of 118 leave-one-out analyses, so it is not driven by one woman, but it remains non-significant after Holm correction throughout.

5. Discussion

5.1 What the Contrast Between Nulls Shows

The same five counts are extreme under population-level exchangeability ($D_M = 5.08$, $p < .001$) and unremarkable under within-woman exchangeability ($D_M = 2.93$, $p = .13$). The global result is real and is not an artefact of the multinomial versus permutation choice (the two agree). It says that vessel events are concentrated in a way a homogeneous population would not produce. The within-woman result says that, conditional on each woman’s own cycle-length values, the order in which they occurred produced no detectable excess of rule-defined events — for these exact statistics, against random reorderings, at the sensitivity documented in Section 4.5. The simplest account consistent with both is between-woman heterogeneity in regularity: women with narrow cycle-length distributions generate many exact repetitions under any ordering, and the global nulls, which pool across women, treat this as excess. This account is supported by the data; it is not uniquely identified by them, and other mechanisms that leave the multiset-conditional distribution of these counts unchanged cannot be excluded.

Two things the H_W result does not show should be stated. It does not show that cycle lengths are serially independent: the autocorrelation documented by Ecochard et al. (2024) in an overlapping cohort is exactly the kind of structure that H_W destroys, and its non-detection here means only that it does not translate into extra exact vessel events in this sample. And it

does not show that the rules are without anticipatory value: whether a woman with three equal consecutive intervals is more likely to repeat the interval is a forecasting question, not addressed here.

5.2 Primary, Exploratory and Robustness Analyses

Primary: the five-pattern joint tests under H_G and H_W , and the Holm-corrected marginal tests under each. *Exploratory:* the Dilug marginal ($p = .085$), the Dilug-in-Dilug deficit under H_W (Holm $p = .19$), the pairwise associations, and the three-pattern subspace. *Robustness:* the multinomial null, the split-batch joint test, truncation, and leave-one-woman-out. No analysis was pre-registered; the five-pattern family was fixed after two rules were found to have zero occurrences (Section 3.1).

5.3 Pattern-Specific Findings

Haflaga is the only pattern that departs from a global null after multiplicity control and it does so strongly (4 of 50,000 replicates reached the observed count). Under H_W its count is below the conditional mean. The 20 women with Haflaga events have markedly lower within-woman SD than women with no events, as the definition of the rule implies; this characterises whom the rule selects rather than validating it. Dilug's global $z = 1.84$ does not survive correction and should be treated as hypothesis-generating; gradual trends (age, weight, hormonal change) would produce arithmetic progressions, but no such mechanism is tested here. The Week pattern in this sample is almost entirely the 28-day cycle (haflaga 29) repeated twice, because longer weekly anchors (35, 42 days) are rare (9.6% of cycles are ≥ 35 days). Week-Dilug events include the 3 Haflaga-30 sequences by containment.

5.4 The Joint Statistic

The quadratic form standardises each component by its null variance and accounts for the null covariance implied by the definitions; Haflaga dominates under the global nulls because its null SD is smallest. Hotelling's T^2 shares the form but assumes multivariate normality of the count vector, which fails for counts as small as 18; the Monte Carlo reference distribution avoids this. A max- T statistic would be less sensitive to simultaneous displacement of several

components; in these data it leads to the same conclusions under all nulls. Using the same replicates to estimate $\hat{\Sigma}$ and to form the reference distribution introduces a finite- B dependence; the split-batch results and the size simulation show it is immaterial at $B = 50,000$. We regard the statistic as a routine adaptation of permutation MANOVA-type tests ([Anderson, 2001](#)) to rule-defined counts and claim nothing further for it.

5.5 Implications

All implications are hypotheses for future study; none is tested here. A woman with an established Haflaga has a within-woman SD below about 1.6 days in this dataset, which is the situation in which calendar-based fertility awareness is most reliable ([Bortot et al., 2010](#)); whether the halachic marker adds to statistical cycle-length prediction is the subject of a companion forecasting study ([Ross, 2026](#)). For observant Jewish women the within-woman result is consistent with the traditional understanding that the rules codify observed regularity; whether the empirical findings bear on the legal standing of the rules is outside the scope of statistics. Changes in vessel status could in principle signal a change in regularity, a clinical hypothesis that would require datasets with clinical endpoints.

5.6 Transferable Lesson

The reusable content of this case is not the statistic but the null-model comparison. Whenever deterministic rules are applied to recurrent-event sequences from many subjects with unequal follow-up — exact inter-seizure intervals in epilepsy diaries, repeated inter-episode intervals in mood disorders, fixed-lag repetitions in behavioural logs — a pooled or i.i.d. null confounds between-subject heterogeneity with within-subject structure, and a within-subject permutation null conditional on each subject’s observed multiset isolates the ordering question at the cost of answering only that question. The two should be reported together, with a simulation showing what the conditional test can detect. Only within-subject data with a defined generative alternative could sharpen this further; a parametric null from a fitted state-space model ([Bortot et al., 2010](#)) would be a natural next step.

5.7 Limitations

Coding and measurement. — Counts are invariant to the $H_k = L_k + 1$ relabelling provided the Week-Dilug anchor is shifted accordingly (verified: 26 Haflaga, 18 Week-Dilug under both conventions; computing on L_k with an unshifted anchor of 30 gives 12 Week-Dilug events). The haflaga distribution declines smoothly through the anchor ($H = 29$: 225 cycles; 30: 175; 31: 141; Figure 11), so the Week-Dilug count does not appear to reflect digit preference; no formal test of heaping was performed. Onset time of day is not recorded, so some events may be misclassified relative to the halachic requirement of a common onset period; a shift that is constant within a woman leaves her multiset and hence the H_W analysis unchanged, but shifts that vary within a woman are not addressed.

Scope. — Calendar dates are absent, precluding the 10 date-dependent rule types; body-sign patterns (*vesset haguf*) are invisible to cycle-length data. The two repeating-cycle rules were dropped after zero occurrences were observed. The sample consists of non-Jewish, regularly cycling, married NFP users aged 18–42 (Fehring et al., 2013a), a self-selected, cycle-aware group who may be more regular than the general population; some halachic authorities hold that regularities in this population need not transfer to observant Jewish women (Ross, 2022b). Results should not be generalised without replication in the target population.

Statistical. — H_W tests exchangeability of each woman’s observed values and nothing more general. The simulations (Section 4.5) show high power against exact-repetition mechanisms and against whole-day autoregression of moderate strength ($\rho \geq 0.5$), but only about 50% power at $\rho = 0.25$ and, necessarily, none against dependence that leaves the distribution of these exact count statistics unchanged; the non-rejection should be read with those limits. The aggregate-count estimand weights women by follow-up; the truncation and leave-one-woman-out analyses indicate that this does not drive the conclusions. Sensitivity to the 18–54 day inclusion window could not be assessed because the excluded cycles are not in the archived file. Pairwise associations involving rare patterns (18 and 27 events) rest on few women and are exploratory.




Figure 11: Frequency of each haflaga value $H_k = L_k + 1$ in the analysed sample. Red bars: $H = 29$ and $H = 31$, the round-number cycle lengths $L = 28$ and $L = 30$ at which digit preference would appear; orange bar: $H = 30$ ($L = 29$), the Week-Dilug anchor. Counts decline smoothly through these values (225, 175, 141), so the Week-Dilug count does not appear to reflect reporting heaping.

6. Conclusion

We applied three exchangeability nulls to the counts of five rule-defined menstrual patterns in 1,554 cycles from 118 women. Under population-level nulls the count vector is far in the tail, driven by a fourfold excess of Haflaga events; under a within-woman permutation null conditional on each woman's own cycle-length values it is not, and all counts lie at or below their conditional means. Simulation confirms that every test used holds its nominal size, that the within-woman tests have high power against exact-repetition alternatives and against autoregressive dependence of moderate strength (about 50% at $\rho = 0.25$, near-certain at $\rho \geq 0.5$), and that the conclusions are insensitive to unequal follow-up. The results are consistent with between-woman heterogeneity in regularity generating the global excess and provide no evidence that within-woman ordering adds rule-defined events; they do not establish the absence of serial dependence, and they say nothing about the forecasting value of the rules, which is the subject of ongoing work. To our knowledge this is the first formal randomisation-based test of the halachic vessel rules on empirical data. The general lesson is that for rule-defined counts in clustered longitudinal sequences the choice of exchangeability null is the analysis, and that a within-subject conditional null should be reported alongside any pooled one.

Ethics Statement

This study is a secondary analysis of a publicly archived dataset (Fehring et al., 2006, 2013a). The original data collection was approved by the Institutional Review Board at Marquette University. No new data were collected, and no individual participant can be identified from the released dataset. Separate ethics approval for secondary analysis was not required under the applicable regulations.

Conflict of Interest

The author declares no conflicts of interest.

Data Availability Statement

The menstrual cycle dataset is publicly available from the Marquette University Natural Family Planning repository (Fehring et al., 2013b). The filtered data file, all analysis code (pattern counting with an equivalence test between two implementations, null generators, simulation validation, post-processing), the Jupyter notebook that produces every figure, and the cached Monte Carlo replicates ($B = 50,000$ per null) that allow every reported number to be reproduced without re-simulation are available at <https://github.com/dvirross/VesetVectorTest>.

AI Assistance Disclosure

Portions of the text were drafted and edited with the assistance of Claude (Anthropic), a large language model, and the tool was also used to generate and refactor $\text{\LaTeX}$ and `Python` code under the author’s direction. All scientific content, analytical decisions and interpretations are the author’s, who takes full responsibility for the accuracy and integrity of the article.

References

- Marti J. Anderson. A new method for non-parametric multivariate analysis of variance. *Austral Ecology*, 26(1):32–46, 2001.
- Thomas B. Berrett, Yi Wang, Rina Foygel Barber, and Richard J. Samworth. The conditional permutation test for independence while controlling for confounders. *Journal of the Royal Statistical Society: Series B*, 82(1):175–197, 2020.
- Paola Bortot, Guido Masarotto, and Bruno Scarpa. Sequential predictions of menstrual cycle lengths. *Biostatistics*, 11(4):741–755, 2010. doi: 10.1093/biostatistics/kxq020.
- Jonathan R. Bull, Simon P. Rowland, Elina Berglund Scherwitsl, Raoul Scherwitsl, Kristina Gemzell Danielsson, and Joyce Harper. Real-world menstrual cycle characteristics of more than 600,000 menstrual cycles. *npj Digital Medicine*, 2:83, 2019.
- Mitchell D. Creinin, Sharon Keverline, and Leslie A. Meyn. How regular is regular? An analysis of menstrual cycle regularity. *Contraception*, 70(4):289–292, 2004.
- Sonya S. Dasharathy, Sunni L. Mumford, Anna Z. Pollack, Neil J. Perkins, Donald R. Mattison, Jean Wactawski-Wende, and Enrique F. Schisterman. Menstrual bleeding patterns among regularly menstruating women. *American Journal of Epidemiology*, 175(6):536–545, 2012.
- Anthony C. Davison and David V. Hinkley. *Bootstrap Methods and Their Application*. Cambridge University Press, Cambridge, 1997.
- René Ecochard, Joseph B. Stanford, Richard J. Fehring, Marie Schneider, Shahpar Najmabadi, and Claude Gronfier. Evidence that the woman’s ovarian cycle is driven by an internal circamonthly timing system. *Science Advances*, 10(15):eadg9646, 2024. doi: 10.1126/sciadv.adg9646.
- Richard J. Fehring, Mary Schneider, and Kathleen Raviele. Variability in the phases of the menstrual cycle. *Journal of Obstetric, Gynecologic & Neonatal Nursing*, 35(3):376–384, 2006.

- Richard J. Fehring, Mary Schneider, and Kathleen Barron. Randomized comparison of two Internet-supported methods of natural family planning. *Contraception*, 88(1):24–30, 2013a. doi: 10.1016/j.contraception.2012.10.010.
- Richard J. Fehring, Mary Schneider, and Kathleen Barron. Natural family planning dataset. Marquette University e-Publications, https://epublications.marquette.edu/data_nfp/7/, 2013b. Dataset accompanying Fehring et al. (2013).
- Sarah Friedrich, Edgar Brunner, and Markus Pauly. Permuting longitudinal data in spite of the dependencies. *Journal of Multivariate Analysis*, 153:255–265, 2017.
- Siôbhán D. Harlow, Xihong Lin, and M. J. Ho. Analysis of menstrual diary data across the reproductive life span: applicability of the bipartite model approach and the importance of within-woman variance. *Journal of Clinical Epidemiology*, 53(7):722–733, 2000.
- Jesse Hemerik and Jelle J. Goeman. Exact testing with random permutations. *TEST*, 27(4): 811–825, 2018.
- Oliver E. Lee and Thomas M. Braun. Permutation tests for random effects in linear mixed models. *Biometrics*, 68(2):486–493, 2012.
- Kirsten Münster, Lene Schmidt, and Peter Helm. Length and variation in the menstrual cycle — a cross-sectional study from a Danish county. *British Journal of Obstetrics and Gynaecology*, 99(5):422–429, 1992.
- Thomas E. Nichols and Andrew P. Holmes. Nonparametric permutation tests for functional neuroimaging: a primer with examples. *Human Brain Mapping*, 15(1):1–25, 2002.
- Fortunato Pesarin and Luigi Salmaso. *Permutation Tests for Complex Data: Theory, Applications and Software*. John Wiley & Sons, Chichester, UK, 2010.
- Belinda Phipson and Gordon K. Smyth. Permutation P-values should never be zero: calculating exact P-values when permutations are randomly drawn. *Statistical Applications in Genetics and Molecular Biology*, 9(1):Article 39, 2010.

- Dvir Ross. Reliability of menstrual cycle reporting in surveys with regard to Jewish religious laws. *Functional Differential Equations*, 26(1–2):1–13, 2019. In Hebrew.
- Dvir Ross. Statistical analysis for the religious laws of abstinence before menstruation. *B.D.D.*, 2021. Accepted 3 September 2021. In Hebrew.
- Dvir Ross. Fixed-difference menstruation Markov chain. *B.D.D.*, 36:77–104, 2022a. In Hebrew.
- Dvir Ross. *Probabilistic and Statistical Analysis of the Menstrual Cycle in a Halachic Context*. PhD thesis, Ariel University in Samaria, Ariel, Israel, 2022b. In Hebrew. Supervisors: Prof. E. Merzbach, Dr. E. Shmerling, Prof. A. Domoshnitsky. Available at <https://github.com/dvirross/PhD>.
- Dvir Ross. Forecasting efficiency of halachic vessel rules: A comparison with statistical cycle-length prediction methods. Manuscript in preparation, 2026.
- Katja M. Schmalenberger, Hafsa A. Tauseef, Jordan C. Barone, Sarah A. Owens, Lynne Lieberman, Marc N. Jarczok, Susan S. Girdler, Jeff Kiesner, Beate Ditzen, and Tory A. Eisenlohr-Moul. How to study the menstrual cycle: Practical tools and recommendations. *Psychoneuroendocrinology*, 123:104895, 2021.
- Laura Symul, Katarzyna Wac, Paula Hillard, and Marcel Salathé. Assessment of menstrual health status and evolution through mobile apps for fertility awareness. *npj Digital Medicine*, 2:64, 2019.
- Gábor J. Székely and Maria L. Rizzo. Energy statistics: A class of statistics based on distances. *Journal of Statistical Planning and Inference*, 143(8):1249–1272, 2013.
- Anderson M. Winkler, Gerard R. Ridgway, Matthew A. Webster, Stephen M. Smith, and Thomas E. Nichols. Permutation inference for the general linear model. *NeuroImage*, 92:381–397, 2014.